

Executive Summary

Vicon Ground-Truth Validation of the AutoPos UWB Self-Localisation System

15 June 2026 | Zekai Xiao

All “§” and “Table” references link to the full report `main_EN.pdf`; numbers are shared with it.

This report presents the first external ground-truth validation of the AutoPos UWB anchor self-localisation system against a Vicon motion-capture reference (28 May 2026, MaD Lab, FAU Erlangen-Nürnberg; §1). Eight DWM1001C-Dev anchors run broadcast SS-TWR with UWB-only inter-anchor self-calibration; a 12-camera Vicon system (sub-millimetre, 0.18 mm median world error) provides truth. The evaluation covers 24 static tag positions, 17 dynamic RotoArm captures, and an offline UWB+IMU fusion replay. All claims are dataset-bound to this measurement.

Key results at a glance

Quantity	Result	Full report
Static 3D accuracy — median / P95 / RMSE	72.7 / 171.5 / 109.8 mm	§5.1, Table 5
Static horizontal / vertical median	37.4 / 61.9 mm	Table 5
Internal repeatability σ_{3D} — unfiltered $\rightarrow$ filtered	67.1 $\rightarrow$ 18.6 mm	Table 6, §6
Anchor self-cal. residual / apparent scale	92.8 mm / +4.36%	§3, §4.2
Dynamic RotoArm floor — track P50, persists	105.8 mm	§7, Table 9
Best UWB+IMU fusion — track P95	231.8 $\rightarrow$ 112.1 mm	§11, Table 18

Headline static accuracy. Under anchor-locked rigid Vicon registration, the production v4-io/T4/F0 pipeline achieves **72.7 mm median 3D error** over 24 static positions (full breakdown in the table above; §5.1, Table 5). This is sub-decimetre median static localisation from a range-only self-calibrated array in this geometry — not a universal UWB specification, since layout, tag placement, vertical aperture, and propagation all affect it.

Anchor self-calibration and the apparent scale bias. The v4-io layout reproduces the Vicon anchor geometry to 92.8 mm median residual (§3, Table 2). The dominant layout-error mode is not random shape distortion but a coherent +4.36% **apparent scale** (mean inter-anchor excess 120.5 mm/pair). A per-anchor decomposition (§4.2, Table 4) traces this to common-mode per-device range bias (mean 94.6 mm/device, all eight terms positive) that the bounded production delay terms under-absorb, so the residual appears geometrically as layout expansion. A re-parameterised solver with a free common mode recovers metric scale to $\approx 1\%$ from UWB ranges alone, leaving sub-percent scale for bench delay calibration.

Key scientific finding — delay–layout coupling. A four-way ablation (layout source $\times$ residual-correction treatment; §4, Table 3) shows that optically surveyed Vicon anchor coordinates *degrade* tag localisation unless the residual delay/range-bias corrections are re-estimated in the Vicon coordinate frame (77.7 mm vs. 108.9 mm RMSE in the ablation estimator). The fitted corrections are **frame-conditioned** and cannot be transplanted across coordinate frames. To our knowledge this is the first explicit four-way ablation showing that externally measured ground-truth coordinates cannot simply replace a self-calibrated UWB layout — the most transferable result of the campaign and the intended core of a methods publication.

Vertical error is systematic, not geometric. The static error is vertically dominated (61.9 mm vertical vs. 37.4 mm horizontal median). A range-only DOP analysis rules out geometric observability: VDOP $\approx$ HDOP over the measured volume (§12, Table 21), and repeatability is *not* vertically

dominated whereas accuracy is, so the deficit is systematic bias, not amplified noise. A height-driver audit identifies the uncorrected **tag-side delay** as the dominant height-dependent lever; the expanded anchor layout partially *masks* it. The 72.7 mm headline therefore benefits in part from this cancellation: holding tag delay at zero on the clean metric-corrected layout gives 109.5 mm median, while an oracle same-hardware tag-delay correction projects a 3D median near 61.7 mm (58.6 mm on the common-mode diagnostic layout). The headline is reported as-measured; these are mechanism-isolation diagnostics, not competing claims.

Two deployment lessons. *Filtering is not calibration* — external position filters cut per-session spread from 67.1 mm to 18.6 mm (§6, Table 8) but leave absolute accuracy essentially unchanged; they remove jitter, not bias. *The dynamic floor is UWB-intrinsic* — RotoArm trajectories show a ~ 100 mm median floor (§7, Table 9) that persists across every layout and tag-solver combination and is not removed by substituting Vicon-measured anchor coordinates, pointing to motion-dependent range bias, antenna orientation, and shadowing rather than layout quality.

UWB + IMU fusion (offline replay). With a synthetic IMU derived from the Vicon antenna-centre trajectory, fusion acts primarily as tail control: the best production row reduces track-level P95 from 231.8 mm to 112.1 mm (§11, Table 18), while UWB absolute positions are required to bound inertial drift. This is replay evidence under datasheet-informed sensor models, not hardware-in-the-loop validation.

Limitations (§13). No hardware time synchronisation (dynamic error is best-fit post-processed, not clock-referenced); a single favourable open-lab site; discrete static positions; constrained circular RotoArm motion; simulated rather than embedded IMU.

Next steps (§13). (1) GPIO/event hardware time-sync; (2) per-device antenna-delay bench calibration plus a one-time or windowed tag-side delay correction, as the dominant vertical-error lever; (3) DOP-aware runtime confidence conditioned on the surviving anchor geometry; (4) closed-loop hardware IMU (ICM-45686 class); (5) unconstrained 3D motion and ≥ 1 cluttered environment; (6) independent ultrasound height-gauge verification. Channel-level diagnostics are also a key open direction: this validation captured ranges only, so the residual error tail and the ~ 100 mm dynamic floor could not be characterised at the channel level; per-link CIR capture (firmware under development) for line-of-sight/non-line-of-sight classification and NLOS-aware range weighting is the planned lever beyond constant per-device delay calibration (see §13).

Bottom line (§14). Within this dataset, AutoPos delivers sub-decimetre median static localisation from a self-calibrated anchor array (up to a manually resolved chirality) in a favourable environment. Two concrete obstacles remain on the path to clinical full-body capture: the apparent scale bias (traced to per-device range bias, largely recoverable by solver re-parameterisation plus delay calibration) and the ~ 100 mm dynamic floor (addressed by time-sync and hardware IMU). Neither is resolved by this dataset.

Acknowledgements. The author thanks Prof. Dr. Björn Eskofier for providing access to the Vicon Motion Capture System at the Machine Learning and Data Analytics Lab (MaD Lab), FAU Erlangen-Nürnberg, and for his continued supervision and guidance throughout this PhD project. The author also thanks Sophie Fleischmann, M. Sc., for coordinating the measurement sessions and for her support with the Vicon system operation during the data collection on 28 May 2026.