

# Novelty Assessment: AutoPos — 3D Anchor Self-Calibration from Inter-Anchor Ranging Alone

System-Paper Positioning and Prior-Art Delineation

Zekai Xiao

June 2026

*This note is the novelty assessment for the AutoPos concept (Paper A, a metrology/deployment systems paper). It is one half of a deliberately split, non-overlapping two-paper program. Paper A reports the working system, its ground-truth-validated accuracy, and the vertical limit it does not cross; the companion mechanism paper (Paper B) analyzes the delay-layout coupling behind that limit. The delay-estimation result appears in Paper A only as an empirical system effect (“modelling per-device delay helps”); Paper B characterizes the coupling itself — its headline results are an anchor-side common-delay $\leftrightarrow$ layout-scale degeneracy and a Vicon-grounded falsification protocol, while a tag-side strand of the same analysis explains why this system’s residual vertical error is a per-tag delay aliasing rather than a geometry deficit that more geometry could remove.*

## 1 One-Sentence Thesis

AutoPos is, to our knowledge, the first UWB anchor self-calibration system that recovers full 3D anchor coordinates jointly with per-device antenna delays from *inter-anchor ranging alone* — with no external sensor (no LiDAR/SLAM/VIO/ultrasound), no mobile survey platform, and no manually surveyed reference positions — and that validates the result against optical ground truth, reporting both the accuracy it reaches and the systematic vertical bias it does not remove. The novelty is a *deployment capability and its honest characterization*, not a new estimator: the individual numerical components are established; what is new is their integration for the inter-anchor-only 3D regime that prior work explicitly avoids, together with ground-truth-free quality diagnostics usable at deployment time.

## 2 Motivation: the Accuracy Ceiling and the Calibration Bottleneck

The systematic review of 37 UWB+IMU motion-tracking systems by [Yogesh et al. \(2023\)](#) documents that even the *best* reported accuracy is only  $\sim 5$  cm: only  $\sim 19\%$  of the systems reach  $\sim 10$  cm even under ideal line-of-sight with optically surveyed anchors, the majority report 10–80 cm, and several exceed 1 m — all well short of the  $\sim 1$  cm that the review derives from matching optical motion capture, the accuracy “gold standard” it considers clinically relevant. Only 5 of the 37 systems attempt 3D positioning at all; the other 32 operate purely in 2D, avoiding the vertical axis. The review singles out systematic UWB ranging error as the dominant bottleneck, and its largest component is per-device antenna delay — a constant ranging bias of  $\sim 150$  mm per node ([Shah et al., 2021](#); [Sang et al., 2019](#)), an order of magnitude above the transceiver’s random noise. The follow-up ([Yogesh et al., 2024](#)) finds that none of the reviewed systems even apply the manufacturer’s recommended

delay-calibration procedure — let alone model the delay inside their estimation pipeline — and reports that even *after* that procedure a  $\sim 3\text{--}4$  cm systematic ranging residual remains.

The systems that do reach high accuracy do so only by importing infrastructure that is unavailable out of the lab: angle-of-arrival base stations and surveyed anchor coordinates (Zihajehzadeh and Park, 2017), or dedicated procedures requiring known reference positions (Kok et al., 2015). This is a *circular dependency*: accurate UWB positioning needs accurately known anchor coordinates, but obtaining those coordinates needs an optical system or surveyed references — precisely what is missing in ambulatory, clinical, and field deployments. Anchor self-calibration is the route out of that circle, and the gap AutoPos targets is that no existing self-calibration method closes it in 3D from inter-anchor ranging alone on validated commodity hardware (Section 3).

**Why this system, and why 4+4: the vertical axis has no shortcut.** The target application is body-worn motion capture, which sharpens the requirement in a way that drone or robot UWB does not. A hovering platform can resolve its height with an onboard downward range sensor (ultrasound, time-of-flight) or fuse a barometer (Xiang et al., 2025); a body-worn tag cannot — an IMU gives only drifting relative vertical displacement, and a barometer only coarse, drifting absolute height, neither adequate for centimetre-level limb tracking. The *only* drift-free, centimetre-level absolute vertical reference available to a body-worn tag is the UWB geometry itself. The vertical must therefore be carried by pure UWB, which is exactly why the array needs vertical geometric diversity (the balanced 4+4 two layers) rather than four coplanar anchors for the horizontal plus a height sensor on the tag. This is also why what we term the *delay-layout coupling* surfaced here at all: forcing UWB to recover its weakest axis (the vertical) is precisely where a common-mode range bias aliases most strongly, so a system required to obtain  $Z$  from ranging alone is the one that exposes the coupling. We name and stake the effect here; the companion paper (Paper B) supplies its identifiability analysis.

### 3 The Gap: A Narrow but Genuinely Unoccupied Combination

UWB self-calibration splits cleanly into two clusters (Table 1), echoing the two-cluster structure of the AutoPos concept’s prior-art survey. The *inter-anchor-only* cluster is restricted to coplanar/2D deployments or requires a few manually surveyed reference anchors; the *3D* cluster achieves height observability only by adding an external sensor (LiDAR, SLAM, or visual-inertial odometry) or a mobile platform that traverses the space — effectively replacing manual surveying with a robot-assisted survey. The combination “3D + inter-anchor ranging only + joint antenna-delay + no known reference” is *unoccupied as a whole*, though every individual axis is occupied — we say so explicitly to avoid an axis-stacked “first.” Joint anchor-plus-delay estimation already exists (Shah et al., 2022; Piavanini et al., 2022; De Preter et al., 2019); 3D inter-anchor-only self-calibration without a known reference exists in simulation (Pereira, 2021) and, via a *mobile* platform on hardware (Batstone et al., 2017). What is new is the *conjunction* on commodity static hardware with optical ground-truth validation, not any one ingredient. The conjunction is genuinely hard, and the field documented it as such and stepped *around* it rather than through it: De Preter et al. (2019) call the 3D calibration “very ill-conditioned” for similar tag/anchor heights and restrict to 2D with tape-measured heights; Corbalan et al. (2023) report vertical resolution “generally poor due to the low spread of anchors along that axis” and recommend measuring height independently; and commercial RTLS deployments that raise a few anchors to recover the vertical *survey* those raised positions rather than self-calibrate them. Each of these resolves the vertical by importing a measured height — the manual step AutoPos removes. AutoPos enters the regime deliberately and reports what actually happens when the vertical is left to ranging alone — which is where the companion paper’s delay finding emerges.

| Method                                  | 3D         | Inter-anchor only      | Delay est. | No known ref. | GT-validated |
|-----------------------------------------|------------|------------------------|------------|---------------|--------------|
| PANS auto-pos. (Decawave/Qorvo, 2019)   | no         | yes                    | no         | yes           | —            |
| Ridolfi 2021 (Ridolfi et al., 2021)     | no (2D)    | yes                    | no         | yes           | 2D           |
| Piavanini 2022 (Piavanini et al., 2022) | no (2D)    | yes                    | yes        | yes           | 2D           |
| Corbalan 2023 (Corbalan et al., 2023)   | no (2D)    | yes                    | no         | <b>no</b>     | yes (2D)     |
| De Preter 2019 (De Preter et al., 2019) | no (2D)    | <b>no</b> (mobile tag) | yes        | <b>no</b>     | RTK-GPS (2D) |
| Kok 2015 (Kok et al., 2015)             | yes        | <b>no</b>              | clock only | <b>no</b>     | yes          |
| Shah 2022 (Shah et al., 2022)           | yes        | <b>no</b>              | yes        | <b>no</b>     | yes          |
| Pereira 2021 (Pereira, 2021)            | yes        | yes                    | no         | yes           | sim. only    |
| Batstone 2017 (Batstone et al., 2017)   | yes        | <b>no</b> (mobile tag) | no         | yes           | real-hw      |
| Nguyen 2025 (Nguyen et al., 2025)       | yes        | <b>no</b> (SLAM)       | yes        | yes           | SLAM-dep.    |
| Yuan 2024 (Yuan et al., 2024)           | yes        | <b>no</b> (LiDAR)      | implicit   | yes           | LiDAR-dep.   |
| Liu & Cao 2025 (Liu and Cao, 2025)      | yes        | <b>no</b> (LiDAR)      | implicit   | yes           | LiDAR-dep.   |
| Delama 2025 (Delama et al., 2025)       | yes        | <b>no</b> (VIO)        | implicit   | yes           | VIO-dep.     |
| <b>AutoPos (this work)</b>              | <b>yes</b> | <b>yes</b>             | <b>yes</b> | <b>yes</b>    | <b>Vicon</b> |

Table 1: Anchor self-calibration prior art. AutoPos occupies the otherwise-empty intersection of 3D, inter-anchor-ranging-only, joint antenna-delay estimation, and no known reference, validated against optical ground truth. “GT-validated” = evaluated against an independent ground-truth positioning system. “Implicit” delay = bias absorbed in optimization rather than explicitly separated.

## 4 What We Contribute

**(1) The capability — and a reference point for later work.** The headline result is a pure-UWB, self-calibrated vertical of  $\approx 62$  mm median absolute error (vs Vicon), near the  $\approx 63$  mm inter-anchor Cramér–Rao floor, from inter-anchor ranging alone — a reference point that later reference-free pure-UWB vertical self-calibration in this deployment regime can be measured against. It is produced by a complete, real-hardware pipeline ( $8\times$  DWM1001C, two elevation levels) that takes bidirectional SS-TWR inter-anchor ranging to a 3D anchor layout with per-device antenna delays, using only commodity single-antenna TWR modules — two orders of magnitude cheaper than the AoA base stations (Zihajehzadeh and Park, 2017) or optical systems prior high-accuracy work relies on, and without the mobile sensor rigs the 3D self-calibration cluster requires. The deployed custom SS-TWR firmware applies a single generic antenna-delay default (the vendor example value, 16436 device units) identically to every node, with no per-device hardware calibration; the per-device delays the solver estimates are precisely the residuals this generic default leaves in the ranging. Calibration is thus pushed entirely into the AutoPos software solve — which is the realistic out-of-the-box deployment condition, and the physical reason a delay term is needed at all.

**(2) The geometry that makes the cell enterable.** A balanced two-layer 4+4 non-coplanar array. Non-coplanarity for 3D is itself known; our claim is narrower and is a statement about the *self-calibration* regime specifically: under reference-free inter-anchor self-calibration — and under the eight-anchor broadcast SS-TWR acquisition the array is dimensioned for — a *balanced* upper layer is what restores well-conditioned vertical *geometry*, whereas an asymmetric raise (e.g. one or two anchors lifted) leaves the vertical ill-conditioned. Deployed systems that raise a few anchors do not hit this wall — but only because they *survey* the raised positions rather than self-calibrate them, sidestepping the conditioning by importing the very measurement AutoPos removes. We quantify this with a self-calibration Fisher analysis (Figure 1): for a single coplanar layer the inter-anchor vertical Cramér–Rao bound is *exactly singular* ( $\partial r/\partial z = 0$ ). A single raised anchor formally unlocks the vertical but leaves it badly conditioned (mean vertical CRLB  $\sim 161$  mm, worst-anchor  $\sim 226$  mm at 4+1); adding upper anchors improves it monotonically. Crucially the *mean* bound saturates early ( $\sim 81$  mm at 4+2,  $\sim 63$  mm at 4+4), so an average alone would suggest stopping at two — but the *worst-conditioned* anchor keeps falling ( $226 \rightarrow 134 \rightarrow 122 \rightarrow 93$  mm across 4+1, 4+2, 4+3, 4+4). The

balanced 4+4 is therefore the smallest array that brings *every* anchor’s vertical to the production floor on average ( $\sim 63$  mm, essentially the deployed median vertical error of  $\approx 62$  mm) and to no worse than  $\sim 93$  mm: the choice over an asymmetric 4+ $k$  raise is a worst-case-uniformity argument, not an observability one. At fixed anchor count a *balanced* (diagonal) pair beats a *clustered* (adjacent) pair by  $1.24\times$ , confirming that it is the balance, not merely the count, that conditions the vertical.




Figure 1: Vertical observability vs. upper-layer anchor count  $k$  (balanced raising). **(A)** Anchor self-calibration vertical CRLB from inter-anchor ranging: singular for a coplanar layer ( $k=0$ ). The *mean* bound (circles) saturates near the  $\sim 63$  mm production vertical error by  $k=2$ , but the *worst-conditioned* anchor (squares) keeps improving to  $k=4$  ( $226 \rightarrow 134 \rightarrow 122 \rightarrow 93$  mm), so the balanced 4+4 is the smallest array with *uniformly* well-conditioned vertical geometry — the case against an asymmetric 4+ $k$  raise. **(B)** Downstream tag-localization DOP for the same idealized balanced array: vertical (VDOP) improves monotonically with  $k$  but remains the worse-determined axis throughout (VDOP > HDOP). The deployment geometry’s measured vertical DOP penalty is reported in the companion mechanism paper.

*Why eight anchors, not more.* The anchor count is set by a system constraint, not chosen for convenience — which answers the natural “why not 8+8 for better geometry?” To keep *moving* body-worn tags consistent, the ranging uses a **broadcast** scheme so that all anchors observe the same transmission instant, eliminating the motion-induced error that sequential per-anchor ranging accumulates as the tag moves; and the system is dimensioned for **10 Hz updates on up to 10 simultaneous tags**. That throughput and motion-coherence budget caps how many anchors can be ranged per cycle, so the balanced 4+4 is the *minimum* two-layer configuration that still reaches the vertical-observability floor (Figure 1) within the budget — a denser array such as 8+8 would either break the 10 Hz $\times$ 10-tag throughput or reintroduce motion error. The claim is therefore not “use more anchors for better geometry” but “reach vertical observability with the fewest anchors a real-time multi-tag broadcast system can afford.”

**(3) Two deployment diagnostics that need no ground truth.** Optical ground truth is exactly what is unavailable at deployment, so we identify two checks that flag a bad calibration using only the UWB data itself.

- *Ranging residual is a misleading quality signal.* With only four anchors — three position unknowns and one redundant measurement — the solver can fit the four distances almost perfectly even when the recovered position is wrong, so the configuration with the *lowest* residual can have the *worst* position. In our data the lowest-residual variant (11.5 mm) gives the worst positioning ( $\sim 169$  mm), while the best positioner (49 mm) has a slightly *higher* residual (12.9 mm). The practical lesson: do not trust ranging residual as a calibration-quality score unless there are redundant anchors ( $\geq 6$ ).
- *Dispersion across dynamic anchor selection exposes hidden bias.* When all 8 anchors are

used for every fix, any systematic calibration error appears as a constant offset and all solver variants look equally good ( $\sim 28\text{--}33$  mm). The differences only appear when the tag selects a *different* 4-anchor subset per fix: a well-calibrated layout still yields consistent positions across subsets, whereas a biased layout yields scattered ones. The spread across subsets is therefore a calibration-quality indicator that a fixed anchor set would hide — and, crucially, it needs no ground truth.

(4) **Honest validation.** We convert the system from a repeatability claim to a ground-truth claim using optical motion capture, and report both the accuracy and the limit (next section).

## 5 What We Are *Not* Claiming (and the honest limits)

- **Not a new estimator.** SDP/MDS initialization, Tukey-bisquare IRLS, MAD/MVUE fusion, and alternating joint antenna-delay estimation are individually established. Joint estimation of anchor coordinates and per-anchor range bias from ranging is itself prior art (De Preter et al., 2019; Piavanini et al., 2022; Shah et al., 2022). Closest on the bias-estimation side is Yogesh et al. (2024), which estimates per-node additive range bias in a fully-connected UWB *swarm* validated against optical ground truth — but with the node positions *known* (from the optical system) and an explicitly pure-additive bias model, so it calibrates bias rather than self-calibrating geometry, and requires the very external reference AutoPos removes. The contribution here is the integration for the 3D inter-anchor-only, reference-free regime and the ground-truth characterization, not the algorithms.
- **The antenna delays are estimated and empirically useful, not yet independently validated as physical delays.** Their correctness is currently inferred only through downstream positioning and scale, so we report them as estimated-and-useful rather than “recovered.” A factory-OTP cross-check was attempted but cannot serve this purpose: across all nine modules the DW1000 OTP stores only a near-uniform lot-level antenna-delay value with no per-device antenna calibration, so it does not supply the per-device delays the solver estimates — a finding that itself matters for the companion paper, where it is reported in full.
- **Consistency is not accuracy.** The preliminary headline numbers (positioning standard deviation of order 5 cm 3D / 4 cm vertical) measure repeatability of repeated fixes, not absolute error. Against optical ground truth the absolute error is larger, and only the ground-truth-validated figures are reported as accuracy.
- **The 4+4 geometry is necessary, not sufficient, for the vertical.** It completes the vertical *geometric* observability, but a systematic vertical bias remains after it — because that residual is not a geometry deficit but a *tag-side* delay aliasing into the vertical. The Fisher analysis makes this concrete: the inter-anchor vertical CRLB saturates near 63 mm for 4+4, at the level of the  $\approx 62$  mm deployed median vertical error — a bound on the *anchor* geometry’s vertical contribution, not on the tag positioning error, but their closeness indicates the geometry is near its floor, so the residual vertical bias is consistent with a tag-side delay alias rather than a geometry deficit. This is the single most important honest caveat, and it is the tag-side face of the delay–layout coupling the companion paper analyzes, distinct from that paper’s *headline* anchor-side delay $\leftrightarrow$ scale result.
- **We do not claim general superiority of self-calibration over surveying** for applications demanding the highest achievable accuracy; we claim a favorable effort/accuracy/cost trade-off for rapid, equipment-free 3D deployment.

## 6 Accuracy by Axis, and the Pure-UWB Vertical

We report two distinct quantities and never conflate them: *repeatability* (standard deviation of repeated static fixes, the headline  $\sim 5$  cm figure) and *absolute accuracy* (error against Vicon, which is what the literature reports). The comparison to the landscape uses the absolute numbers.

**Horizontal (X–Y plane).** Repeatability  $\sim 27$  mm; absolute  $\sim 37$  mm (median vs Vicon). This sits in the same band as the horizontal accuracy of the best fusion systems ( $\sim 30$ – $40$  mm per axis). The horizontal is the easy axis for everyone; we neither lead nor lag here.

**Vertical — the standout.** Repeatability  $\sim 41$  mm; absolute  $\sim 62$  mm (median vs Vicon), from *pure UWB* with self-calibrated anchors, no IMU fusion, no AoA hardware, no surveyed anchors, and no external height sensor. This is the axis the field cannot get cheaply: 32 of the 37 reviewed systems avoid it (2D); the systems that do resolve it import a crutch a body-worn tag cannot use — AoA base stations plus biomechanical constraints plus OptiTrack (Zihajehzadeh and Park, 2017), or surveyed anchors plus IMU fusion. The pure-UWB vertical here lands in the band that others reach only with external help. **The defensible claim is not “smallest vertical number” but “a usable vertical recovered from UWB alone, under the body-worn constraint that rules out every external height aid”** — which is precisely why the balanced 4+4 array exists.

**3D total.** Repeatability  $\sim 50$  mm ( $\sim 28$  mm when all eight anchors are visible online); absolute  $\sim 73$  mm median (P95  $\approx 172$  mm, RMSE  $\approx 110$  mm). This sits in the upper range of the reviewed systems (10–80 cm) and within reach of the best ( $\sim 48$  mm 3D RMSE), though we do *not* claim it is the single smallest absolute number: the best reviewed result ( $\sim 48$  mm, a 24 s walk with surveyed anchors and IMU fusion) still edges it. The  $\approx 73$  mm median further benefits, in part, from the self-calibrated production layout’s expanded scale compensating an uncorrected tag-side delay: with the tag-side delay left at zero, the clean metric-corrected layout gives a larger  $\approx 110$  mm median (the mechanism and its in-environment control are analyzed in the companion paper).

**Honest caveat on the comparison.** The reviewed systems are UWB+IMU *fusion* tracking a *moving* tag; AutoPos is UWB-only positioning of a *static* tag. Fusion and IMU help them; the static setting helps us. So the correct statement is that AutoPos reaches the same accuracy band *UWB-only and self-calibrated, with a resolved vertical*, not that it beats fusion systems.

**Reproducibility scope — the acquisition scheme is part of the method.** The positioning figures above are obtained at the system’s designed operating point: an eight-anchor, two-layer **broadcast SS-TWR** acquisition in which every fix draws on the full anchor set. This is constitutive, not a favourable selection. A reproduction that drives the same layout with a naive minimal-anchor SS-TWR scheme, or that collapses the array to a single coplanar layer, will *not* recover these numbers — the vertical in particular returns to the ill-conditioned  $\partial r / \partial z \rightarrow 0$  regime the Fisher analysis proves exactly singular for a coplanar layer. The graceful, layer-asymmetric degradation under partial anchor dropout (losing an *upper* versus a *lower* anchor is not equivalent) is characterised by the stratified keep- $k$  analysis in the full system evaluation.

## 7 Relationship to the Companion Mechanism Paper

The companion mechanism paper (Paper B) is the natural home for the finding that this system surfaces: that the residual vertical error is the per-tag projection of a common-mode range bias — the UWB analogue of the GPS receiver-clock/altitude coupling — which no amount of added geometric diversity removes, only an explicit tag-delay correction (demonstrated there by an active vertical-injection control that produced no dose-response). Keeping the two separate avoids cannibalizing the higher-tier mechanism paper with the systems paper, and lets the mechanism paper cite this validated system as the substrate for its analysis. Paper A establishes priority on the deployment capability and the balanced 4+4 geometry; Paper B establishes priority on the delay–layout coupling — headlined by the anchor-side delay $\leftrightarrow$ scale identifiability analysis and the falsification protocol,

with the tag-side vertical aliasing above as the strand that bears on this system’s residual limit.

## 8 Target Venue

A systems/deployment metrology venue: *IEEE Sensors*, *IEEE Access*, or *IPIN* (full paper). These reward a complete, ground-truth-validated, deployable system with a thorough ablation; the individually-known components are acceptable there given the integration and evaluation. The higher-tier *Measurement / IEEE TIM* target is reserved for the companion mechanism paper. Submitting Paper A first establishes the system as the cited substrate for Paper B.

## 9 Risk Assessment

- **“Incremental / just a combination.”** Mitigated by leading with the empty-cell capability and the ground-truth-free diagnostics rather than the (obvious) two-layer idea, and by the honest validation. The hook is “first ground-truth-validated 3D inter-anchor-only self-calibration + deployment diagnostics,” not “we used two layers.”
- **“Does it actually work in 3D?”** Pre-empted by reporting the validated accuracy and the residual vertical limit openly, and by pointing to the companion mechanism paper rather than over-claiming.
- **“Two layers trivially fix Corbalan’s Z problem.”** Answered by the companion’s finding: balanced 4+4 fixes the *geometric* part, but the vertical bias persists for a *delay* reason — so the contribution is not the layout but the demonstration that the obvious geometric fix is necessary-but-insufficient.
- **Prior-art overlap.** De Preter, Piavanini, Corbalan, Shah are cited prominently and delineated on the 3D / inter-anchor-only / no-known-reference axes; the overlap on *method* is acknowledged, and the contribution is scoped to capability and characterization.

## 10 Publication Strategy and Recommended Manuscript Front-Matter

**Priority strategy.** The most defensible novel element here is the balanced 4+4 array for *pure-UWB vertical observability* in reference-free self-calibration. To establish it as citable prior work, the recommended path is to keep the two papers separate and publish this system paper *first and early* (a public, dated, indexed record — e.g. an arXiv preprint followed by a fast systems venue), with the companion mechanism paper following and citing it. A merged paper would subordinate the 4+4 result to a delay-coupling narrative, reducing its discoverability and delaying the priority record. (The arXiv/early-disclosure decision interacts with institutional IP policy and patentability and should be agreed with the supervisor before posting.)

**Enabling disclosure (for defensive-publication value).** To maximise the blocking value of the priority record against a later overlapping claim, the manuscript states the full combination once, as a single enabling configuration: 3D anchor self-calibration; inter-anchor ranging only; joint per-device antenna-delay estimation; commodity single-antenna DWM1001C modules; a balanced two-layer 4+4 array; and optical-ground-truth validation — with no survey, no external *Z* sensor, and no mobile platform. A dated, indexed disclosure of this exact conjunction is what makes the combination both unpatentable by others and unambiguous to cite as prior art.

**Reproducibility (a citation magnet, by design).** The Fisher/CRLB derivation and figure ship as a runnable artifact, and the inter-anchor ranging dataset with solver outputs accompanies the

preprint, so the  $\approx 63$  mm vertical floor and the by-axis accuracy can be reproduced independently. Releasing the derivation and data is part of the priority strategy: a reusable artifact that later work must cite to build on does more to lock in prior-art status than the claim sentence alone.

**Recommended title** (claim-forward, for discoverability): “*Pure-UWB Vertical Observability from a Balanced Two-Layer (4+4) Anchor Array: Reference-Free 3D Anchor Self-Calibration with Joint Antenna-Delay Estimation.*” Alternatives retain the “AutoPos” name or lead with “Recovering the Vertical from UWB Alone.”

**Recommended abstract.**

Anchor-based ultra-wideband (UWB) positioning requires accurately known anchor coordinates, yet for body-worn motion capture the vertical axis has no shortcut: unlike a hovering platform, a body-worn tag cannot carry a downward range sensor or usefully fuse a barometer, so a drift-free, centimetre-level vertical reference can come only from the UWB geometry itself. We present a reference-free 3D anchor self-calibration that recovers anchor coordinates and per-device antenna delays from inter-anchor two-way ranging alone — no external sensor, no mobile survey platform, no surveyed reference positions — on commodity (DWM1001C) hardware, validated against optical (Vicon) ground truth. The enabling element is a balanced two-layer (4+4) anchor array: via a Fisher/Cramér–Rao analysis we show a single coplanar layer leaves the vertical exactly unobservable ( $\partial r / \partial z = 0$ ), that raising anchors unlocks it, and that the balanced 4+4 array reaches a vertical CRLB floor of  $\approx 63$  mm, with *balance* — not merely anchor count — conditioning the vertical; the eight-anchor count itself follows from a real-time broadcast protocol dimensioned for 10 Hz on up to 10 simultaneous tags. Reporting by axis and separating repeatability from absolute accuracy, the system attains  $\approx 5$  cm 3D /  $\approx 4$  cm vertical repeatability and  $\approx 7$  cm 3D /  $\approx 6$  cm absolute (vs Vicon) — a pure-UWB vertical where most comparable systems avoid the axis or reach it only with angle-of-arrival, biomechanical, or surveyed-anchor aids. The 4+4 geometry is necessary but not sufficient: a residual systematic vertical bias remains, attributed to a tag-side delay aliasing into the vertical (companion paper). To our knowledge this is the first demonstration that the *vertical* axis of a UWB anchor array can be self-calibrated to the centimetre level —  $\approx 62$  mm median absolute against optical ground truth, near a  $\approx 63$  mm inter-anchor Cramér–Rao floor — from inter-anchor ranging alone on commodity single-antenna modules, with the delay coupling that bounds it identified and quantified.

**Recommended keywords.** ultra-wideband (UWB); anchor self-calibration; anchor auto-positioning; reference-free / infrastructure-free localization; inter-anchor two-way ranging (TWR); two-layer anchor array; vertical observability; 3D positioning; antenna-delay calibration; body-worn / wearable motion capture; Cramér–Rao bound; dilution of precision.

**Citable priority sentence** (to appear verbatim in the body): “*A balanced two-layer (4+4) anchor array renders the vertical well-conditioned for reference-free, inter-anchor-only 3D UWB self-calibration — a single coplanar layer is exactly singular ( $\partial r / \partial z = 0$ ) and a single raised anchor only weakly unlocks it, while the balanced 4+4 reaches a  $\approx 63$  mm vertical Cramér–Rao floor — enabling a pure-UWB vertical without any external height aid.*”

## References

- Batstone, K., Oskarsson, M. and Åström, K., 2017. Towards real-time time-of-arrival self-calibration using ultra-wideband anchors. In: *International Conference on Indoor Positioning and Indoor Navigation (IPIN)*, pp.1–8.
- Corbalan, P., Picco, G.P., Coors, M. and Jain, V., 2023. Self-localization of ultra-wideband anchors: From theory to practice. *IEEE Access*, 11, pp.29711–29725.
- De Preter, A., Goysens, G., Anthonis, J., Swevers, J. and Pipeleers, G., 2019. Range bias modeling and autocalibration of a UWB positioning system. In: *International Conference on Indoor Positioning and Indoor Navigation (IPIN)*, pp.1–8.



- Decawave Ltd., 2019. *MDEK1001 Kit User Manual* (DWM1001 module development and evaluation kit; PANS auto-positioning). Version 1.2.
- Delama, G., Borowski, I., Jung, R. and Weiss, S., 2025. Real-time initialization of unknown anchors for UWB-aided navigation. In: *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. arXiv:2506.15518.
- Kok, M., Hol, J.D. and Schön, T.B., 2015. Indoor positioning using ultrawideband and inertial measurements. *IEEE Transactions on Vehicular Technology*, 64(4), pp.1293–1303.
- Nguyen, T.-D., Nguyen, T.-M. and Nguyen, V.-H., 2025. Coordinate-consistent localization via continuous-time calibration and fusion of UWB and SLAM observations. *arXiv preprint* arXiv:2510.05992.
- Liu, X. and Cao, M., 2025. Robust simultaneous UWB-anchor calibration and robot localization for emergency situations. *arXiv preprint* arXiv:2503.22272.
- Pereira, J., 2021. *Anchor self-calibrating schemes for UWB based indoor localization*. M.Sc. thesis, Rochester Institute of Technology.
- Piavanini, M., Barbieri, L., Brambilla, M., Cerutti, M., Ercoli, S., Agili, A. and Nicoli, M., 2022. A self-calibrating localization solution for sport applications with UWB technology. *Sensors*, 22(23), 9363.
- Ridolfi, M., Fontaine, J., Van Herbruggen, B., Joseph, W., Hoebeke, J. and De Poorter, E., 2021. UWB anchor nodes self-calibration in NLOS conditions: A machine learning and adaptive PHY error correction approach. *Wireless Networks*, 27, pp.3007–3023.
- Sang, C.L., Adams, M., Hörmann, T., Hesse, M., Pörmann, M. and Rückert, U., 2019. Numerical and experimental evaluation of error estimation for two-way ranging methods. *Sensors*, 19(3), 616.
- Shah, S., Chaiwong, K., Kovavisaruch, L., Kaemarungsi, K. and Demeechai, T., 2021. Antenna delay calibration of UWB nodes. *IEEE Access*, 9, pp.63294–63305.
- Shah, S., Kovavisaruch, L., Kaemarungsi, K. and Demeechai, T., 2022. Node calibration in UWB-based RTLSs using multiple simultaneous ranging. *Sensors*, 22(3), 864.
- Xiang, Z., Chen, L., Wu, Q., Yang, J., Dai, X. and Xie, X., 2025. An improved UWB indoor positioning approach for UAVs based on the dual-anchor model. *Sensors*, 25(4), 1052.
- Yogesh, V., Buurke, J.H., Veltink, P.H. and Baten, C.T.M., 2023. Integrated UWB/MIMU sensor system for position estimation towards an accurate analysis of human movement: A technical review. *Sensors*, 23(16), 7277.
- Yogesh, V., Grevinga, L., Voort, C., Buurke, J.H., Veltink, P.H. and Baten, C.T.M., 2024. Novel calibration method for improved UWB sensor distance measurement in the context of application for 3D analysis of human movement. *Engineering Science and Technology, an International Journal*, 58, 101844.
- Yuan, S., Lou, B., Nguyen, T.-M., Yin, P., Cao, M., Xu, X., Li, J., Xu, J., Chen, S. and Xie, L., 2024. Large-scale UWB anchor calibration and one-shot localization using Gaussian process. *arXiv preprint* arXiv:2412.16880.
- Zihajehzadeh, S. and Park, E.J., 2017. A novel biomechanical model-aided IMU/UWB fusion for magnetometer-free lower body motion capture. *IEEE Transactions on Systems, Man, and Cybernetics: Systems*, 47(6), pp.927–938.